



Final Presentation: Keyboard Snooping

An Le, Eugene Chu, Zixuan Zhong



Overall Project Goals and Specific Aims


- Perform keyboard snooping with microphone and smartwatch (motion sensors)
- Given audio and motion sensor (accelerometer, gyroscope, or linear accelerometer) data
 - Identify specific keys
 - Identify 6-digit PINs
 - Identify tasks involving keyboard

Technical Approach

- Algorithms
 - Key identification
 - Audio: lossless wave → extract peaks → FFT → NN classifiers
 - Motion Sensor: extract displacement manually? direct training?
 - Task identification
 - Neural network classifier
- Datasets
 - Task 1: single keys, each has an audio wave file
 - Task 2: 6-digit PINs, each has an audio wave file and five accelerometer data
 - Task 3: each sample is at least 30 seconds long and has both audio data and accelerometer data
- Platform
 - Data collection:
 - Audio Recording: A Python Program (software), The built-in microphone of macbook pro (hardware)
 - Motion Sensor: Android Physical Toolbox Suite (app), any android phone (simulating a smart watch)
 - Training and inference:
 - Keras

Task 1 Implementation: Recording and Keylogging

- A python program instruct the user to type
 - Audio: Macbook built-in stereo microphone (.wav)
 - Keylogger: Python **keyboard** module (.csv)



Current devices:

```
> 0 Built-in Microphone, Core Audio (2 in, 0 out)
< 1 Built-in Output, Core Audio (0 in, 2 out)
```

```
Choose mic (index): 0
Recording parameters set.
Started microphone . . .
Please start actual recording by pressing [SPACE] . . .
```

Start actual recording:


```
1) space 2 enter
2) enter 3 enter
3) enter 4 enter
4) enter 5 enter
5) enter 6 enter
```

Stop recording.

```
Writing key logs to outputs/0320074758.csv . . .
```

Task 1 Implementation: Preprocessing and Architecture

- Audio processing
 - Peak extraction
 - Detect peaks given threshold and minimum time between keystrokes
 - Ensure threshold large enough to filter out release peaks
 - Grab audio from 5ms before peak to 20ms after peak
 - Feature extraction
 - Sliding window of length 20ms and stride of 1ms
 - Take average of FFTs of sliding windows
 - Each feature normalized
- Neural network
 - MLP with two hidden FC layers (64, 64), dropout (0.5), FC softmax output (36)
 - Output layer size 36 (a-z, 0-9)



Task 1 Implementation: Machine Learning and Result

- Dataset
 - 4807 total keystrokes (a-z, 0-9)
 - Train-val-test split: 70%-15%-15%
- Training
 - Batch size 128
 - 100 epochs
 - Save model with lowest validation loss
- Results (top-1 accuracy)
 - Single model: **96-98%**
 - Ensemble (averaging) of 5 models: **99%**
 - cf. Keyboard Acoustic Emanations (2004): 79%


Task 2 Implementation:

Which sensor should we use and how?

- Can we extract displacement manually?
 - Conversion from the device's frame to the world's frame, and Double integration from accelerometer to displacement (high school physics)
 - Need linear accelerometer data and orientation data (low sampling rate)
 - Large error even the Kalman Filter is used
 - We lose details when doing manual extraction
 - Previously: a sequence of sensor data; After: only one value
- Which motion sensor to use, accelerometer, gyroscope, or linear accelerometer?
 - Trained three model (accelerometer, gyroscope, or linear accelerometer)
 - Gyroscope: hard to find peaks
 - Accelerometer V.S. linear accelerometer: 80% > 70%

Task 2 Implementation: Hardware and Software

- ZenWatch 2 & AndroidSensorDemo
 - NESL has a app for ZenWatch 2
 - <https://github.com/nsl/AndroidSensorDemo>
 - which can dump GRAVITY, ACCELEROMETER, and LINEAR_ACCELEROMETER sensor data
 - Sampling rate not high enough
 - Some logging issues
- Android phone & Physics Toolbox Sensor Suite
 - High sampling rate: 400 Hz
 - All sensors
 - Single sensor or customized combinations of multiple sensors



Task 2 Implementation: Synchronization and Labeling

Synchronization

1. Find the peak from the first X seconds of the audio data
2. Find the peak from the first Y seconds of the accelerometer data
3. Based on keylogs, cut a recording into groups
 - a. Each group contains a single key (Task 1)
 - b. Each group contains 6 digits (Task 2)
4. For the PIN task, Find 6 peaks for each group and cut the accelerometer data into 5 parts

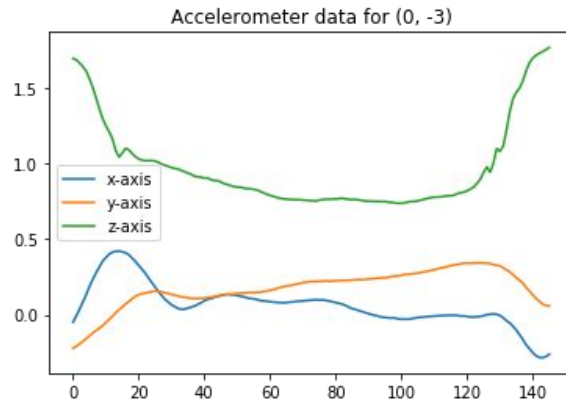
Labeling

- Labels of audio data: from key logs
- Labels of accelerometer data: equivalent displacements
 - Encode each key to 2D coordinates: 0: (0, 0), 1: (-1, 1), 2: (0, 1), 3: (1, 1), 4: (-1, 2), 5: (0, 2), 6: (1, 2), 7: (-1, 3), 8: (0, 3), 9: (1, 3)
 - If the PIN is "104582", there are 5 key pairs in this PIN, which are: <1, 0>, <0, 4>, <4, 5>, <5, 8>, <8, 2>
 - Each key pair is labeled with a vector: (1, -1), (-1, 2), (1, 0), (0, 1), (0, -2)

Task 2 Implementation:

Preprocessing and Architecture

- Audio processing
 - Same as Task 1
- Accelerometer data processing
 - Use segments between push peaks
 - Tried FFT but accuracy low
 - Instead, we resample segments to fixed length of 256
- Neural network
 - Audio
 - Same as Task 1 except with 10 output classes (0-9)
 - Accelerometer
 - Similar architecture to audio
 - Linear output layer representing 2D coordinates
 - MSE loss function



Task 2 Implementation: Machine Learning and Result

- Dataset
 - 718 PINs (0-9) of length 6
 - 2 people, 2 keyboards
 - For each PIN, 6 audio samples and 5 accelerometer samples
 - Train-val-test split: 70%-15%-15%
- Training
 - Batch size 128 for both models
 - 100 epochs for both models
 - Save models with lowest validation loss
- Key model evaluation
 - Individual key accuracy: **84.1%**
 - Top-3 accuracy: **95.1%**
- Displacement model evaluation
 - Accuracy (after rounding output): **80.8%**
 - Loss (MSE): **0.172**

Task 2 Implementation:

Maximum Likelihood Tree Search (MLTS)

- Main idea: find the most probable PIN given key probabilities, displacement estimates assuming that
 - Length of PIN is known
 - Distribution of displacement estimate is bivariate normal with $\sigma_x = \sigma_y$
- Objective function (log likelihood)
 - **Every subterm is effectively non-positive**

$$LL(G) = \sum_{i=1}^{|G|} \log(p(k_i = g_i)) - c \sum_{i=2}^{|G|} |d_i - \hat{d}_i|^2$$

G = PIN subguess with g_i as the i^{th} key

K = True PIN with k_i as the i^{th} key

p = Softmax outputs of key estimator

c = Reliability constant of displacement estimator

d_i = True displacement from g_{i-1} to g_i

\hat{d}_i = Estimate of above

Task 2 Implementation:

Maximum Likelihood Tree Search (MLTS)

- How to find maximum likelihood
 - Naive search: Try all PIN combinations (1 million!)
 - **With heuristic: Set baseline to LL of 6-key PIN guess from key classifier**
 - Perform DFS starting with 1-key subguesses
 - If $LL(\text{subguess})$ less than or equal to maximum LL so far, **do not go deeper**
 - Otherwise, add one more key to subguess
 - Update maximum LL if subguess has 6 keys and its LL is greater than the current max
- **Example run**
 - Set baseline
 - Key classifier guesses [6, 9, 5, 0, 9, 9]
 - $LL([6, 9, 5, 0, 9, 9]) = -7.424$
 - Prune bad guesses
 - $LL([1]) = -7.688 \leq -7.424$
 - Skip all guesses starting with 1
 - Update max LL and best guess if better full guess is found
 - $LL([8, 9, 5, 0, 9, 9]) = -5.963 > -7.424$
 - Imagine $d1 = [1, 0]$
 - Turns out the true PIN is [8, 9, 5, 0, 9, 9]

Task 2 Implementation: Evaluation after Audio and Accelerometer Fusion

- Individual keys
 - Key classifier only: **84.1%**
 - With fusion + MLTS: **91.8%**
- Entire PIN (5 attempts)
 - Key classifier only:
 - With fusion + MLTS: **83.3%**
- **Additional information from accelerometer reduces error rate**

Task 2 Implementation: Evaluation on Additional Test Set

- Different person, different keyboard
- Displacement estimator
 - Accuracy (after rounding output): **51.8%**
 - Loss (MSE): **0.420**
- Key classifier
 - Key classifier only: **13.6%**
 - With fusion + MLTS, $c = 1$: **13.6%**
 - With fusion + MLTS, $c = 100$: **33.3%**
- Performance drop of displacement estimator not as drastic as that of key classifier since hand movement is fairly consistent across different people
- **Some information gain from displacement despite poor key classification**

Task 3 Implementation: Hardware and Software

In this task we collect sensor fusion data from three user tasks (gaming, messaging, emailing) for classification


- Same setup as Task 2 (Android phone & Physics Toolbox Sensor Suite)
- Dataset: 30 seconds/sample, 30 samples/task
- Main idea: find the most probable task given, accelerometer and audio information

Task 3 Implementation: Preprocessing and Architecture

- For this task, we preprocess accelerometer and audio data to reduce feature size:
 - Calculate mean and standard deviation of acceleration for every 300 samples
 - Fourier Transform on audio waveform with window size of 20ms and stride of 1ms
 - Final feature size: 631 per sample
 - Train-val split: 85%-15%
 - Batch size: 10
 - Epochs: 100
- Neural network
 - Audio
 - FFT sample features (1D)
 - Accelerometer
 - Flatten input for FC layer
 - FC layer with batchnorm and dropout
 - Categorical cross-entropy

Task 3 Implementation: Machine Learning and Result

- Results:
 - High accuracy, but may be have overfitted
 - Model easily converges because of small dataset
- Task classification is also dependent on the key press intervals (Intuition):
 - Gaming tasks have constant key presses limited to certain keys (little acceleration change)
 - Messaging has multiple keys pressed clustered in short intervals (requires the other person to reply)
 - Emailing tasks are have more constant typing over a longer period (longer sentences)
 - Dataset is too small



Prior work examples and Relative Novelty:

Task 2

- Keyboard Acoustic Emanations (D. Asonov et al., 2004)
 - First work to develop key classifier based on sound of keypress
 - Works because **the keyboard plate is struck at different spots, producing different timbres like a drum**
- Keyboard Acoustic Emanations Revisited (L. Zhuang et al., 2005)
 - Utilizes clustering and language model to predict keys and words without prior labeling
 - Analyzes key identification attack as substitution cipher
- Don't Skype & Type! Acoustic Eavesdropping in Voice-Over-IP (A. Compagno et al., 2017)
 - Considers different attack scenarios such as knowing or not knowing actual keystrokes or laptop manufacturer
 - Relies on a database of sounds from different keyboards to infer victim's laptop
- **Our work: Fusion of audio and motion sensor data**
 - Exploits estimated displacement from accelerometer to improve key prediction
 - No language model required

Prior work examples and Relative Novelty:

Task 3

- Activity classification using a single wrist-worn accelerometer (S. Chernbum roong et al., 2011)
 - Two models: “Decision Tree C4.5” and “1-layer FC network” for classification of 5 tasks
 - sitting, standing, lying, walking, and running
 - Using **Correlation-based Feature Selection** gives the best accuracy of 94.13% and F-1 score of 0.91
- A Comprehensive Study of Activity Recognition Using Accelerometers (N. Twomey et al. 2018)
 - The paper surveys and evaluates methods of activity recognition using accelerometers.
 - A window length of 5 to 7 seconds gave the highest prediction accuracy using a single wristband accelerometer.
- **Our Work: Fusion of audio and motion sensor data to infer tasks with relatively similar movements**

Strengths, Weakness, and Future Directions

- Strength
 - Multiple scenarios (3 tasks)
 - Multiple ways of sensor fusion
 - High accuracy
- Weakness
 - Applicability. Samples from limited number of hardware.
 - The model requires high sampling rate which is hard to be achieved
 - The Victim may not use the hand with watch to type
- Future Directions
 - More data samples
 - Data samples from different hardware devices
 - To identity the user's emotion while typing

Contributions

An Le	Survey of Prior Work, Machine Learning (Task 1 & 2), Maximum Likelihood Search Algorithm (Task 2)
Eugene Chu	Dataset Generation (Task 1 & 2 & 3), Machine Learning (Task 3)
Zixuan Zhong	Explorations on Motion Sensors (Frame Conversion, Double Integration, and Sensor Selection), Dataset Generation (Task 1 & 2 & 3), and Dataset Synchronization (Task 1 & 2 & 3)

GitHub Repo and Project Website

- Github Repo
 - <https://github.com/binhanle/keyboard-snooping>
- Project Website
 - <https://binhanle.github.io/keyboard-snooping/>

References

- Android Physical Toolbox Suite: <https://www.vieyrasoftware.net>
- Python3: <https://www.python.org>
- Keyboard: <https://github.com/boppreh/keyboard>
- Keras: <https://keras.io>